

```
In [ ]: import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import numpy as np

/tmp/ipykernel_48614/1109543917.py:1: DeprecationWarning:
Pyarrow will become a required dependency of pandas in the next major release of pandas (pandas 3.0),
(to allow more performant data types, such as the Arrow string type, and better interoperability with other libraries)
but was not found to be installed on your system.
If this would cause problems for you,
please provide us feedback at https://github.com/pandas-dev/pandas/issues/54466

import pandas as pd
```

EDA

```
In [ ]: df = pd.read_csv("dataset/scaler_clustering.csv")
df.head()
```

Out[]:

	Unnamed: 0	company_hash	email_hash	orgyear	ctc	job_position	ctc_updated_year
0	0	atrgxnnt xzaxv	6de0a4417d18ab14334c3f43397fc13b30c35149d70c05...	2016.0	1100000	Other	2020.0
1	1	qtrxvzwt xzegwgb rxbxnta	b0aaf1ac138b53cb6e039ba2c3d6604a250d02d5145c10...	2018.0	449999	FullStack Engineer	2019.0
2	2	oizwnvwnxw vx	4860c670bcd48fb96c02a4b0ae3608ae6fdd98176112e9...	2015.0	2000000	Backend Engineer	2020.0
3	3	ngpggutaxv	effdede7a2e7c2af664c8a31d9346385016128d66bbc58...	2017.0	700000	Backend Engineer	2019.0
4	4	qxen sqghu	6ff54e709262f55cb999a1c1db8436cb2055d8f79ab520...	2017.0	1400000	FullStack Engineer	2019.0

```
In [ ]: df.shape
```

Out[]: (205843, 7)

```
In [ ]: df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 205843 entries, 0 to 205842
Data columns (total 7 columns):
#   Column          Non-Null Count  Dtype
---  ---
0   Unnamed: 0      205843 non-null  int64
1   company_hash    205799 non-null  object
2   email_hash      205843 non-null  object
3   orgyear         205757 non-null  float64
4   ctc             205843 non-null  int64
5   job_position    153279 non-null  object
6   ctc_updated_year 205843 non-null  float64
dtypes: float64(2), int64(2), object(3)
memory usage: 11.0+ MB
```

```
In [ ]: df.drop(columns=["Unnamed: 0"], axis=1, inplace=True)
```

```
In [ ]: df.describe()
```

Out[]:

	orgyear	ctc	ctc_updated_year
count	205757.000000	2.058430e+05	205843.000000
mean	2014.882750	2.271685e+06	2019.628231
std	63.571115	1.180091e+07	1.325104
min	0.000000	2.000000e+00	2015.000000
25%	2013.000000	5.300000e+05	2019.000000
50%	2016.000000	9.500000e+05	2020.000000
75%	2018.000000	1.700000e+06	2021.000000
max	20165.000000	1.000150e+09	2021.000000

```
In [ ]: df['orgyear'] = df['orgyear'].astype("object")
df["ctc_updated_year"] = df["ctc_updated_year"].astype("object")
```

```
In [ ]: df.describe()
```

Out[]:

	ctc
count	2.058430e+05
mean	2.271685e+06
std	1.180091e+07
min	2.000000e+00
25%	5.300000e+05
50%	9.500000e+05
75%	1.700000e+06
max	1.000150e+09

```
In [ ]: df.describe(include="object")
```


Out[]:

	company_hash	email_hash	orgyear	job_position	ctc_updated_year
count	205799	205843	205757.0	153279	205843.0
unique	37299	153443	77.0	1016	7.0
top	nnnv wgzohrmvzwj otqcxwto	bbace3cc586400bbc65765bc6a16b77d8913836cfc98b7...	2018.0	Backend Engineer	2019.0
freq	8337	10	25256.0	43554	68688.0

```
In [ ]: fig=plt.figure(figsize=(10,4)) # width*height

# Create a KDE plot with filtered data
plt.subplot(1,1,1)
sns.kdeplot(df["ctc"], fill=True)
```

Out[]: <Axes: xlabel='ctc', ylabel='Density'>



- CTC seems to have outliers, we will be handling the same while working further.
- note: outliers doesn't mean that they are noise, in case of income we may have people with expemtional income.

```
In [ ]: fig=plt.figure(figsize=(10,4)) # width*height

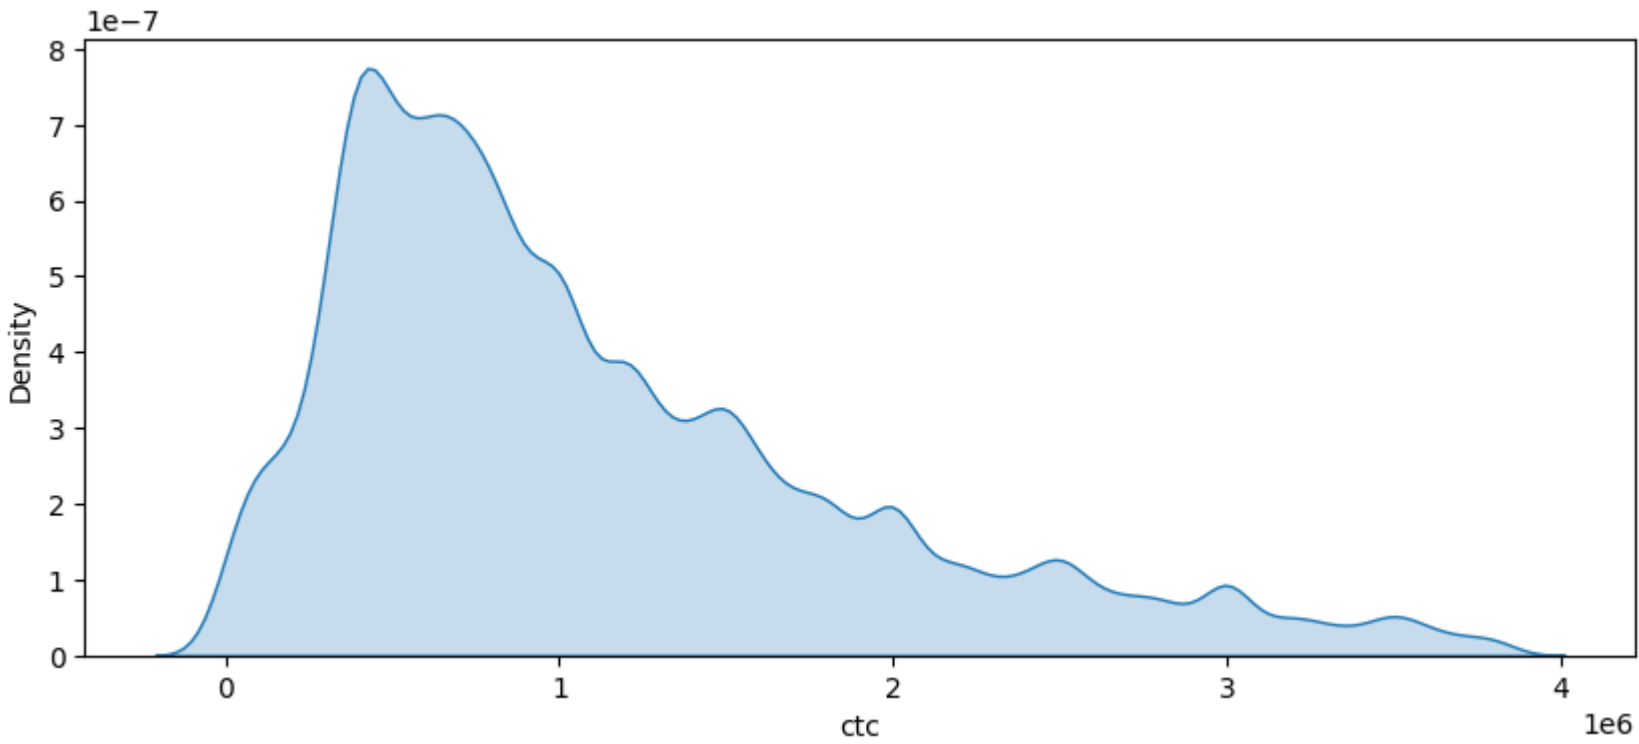
# Get the data
data = df['ctc']

# Set a range to exclude outliers (adjust as needed)
lower_limit = np.percentile(data, 0)
upper_limit = np.percentile(data, 95)

# Filter the data to exclude outliers
filtered_data = data[(data >= lower_limit) & (data <= upper_limit)]

# Create a KDE plot with filtered data
plt.subplot(1,1,1)
sns.kdeplot(filtered_data, fill=True)
```

Out[]: <Axes: xlabel='ctc', ylabel='Density'>



- From the CTC marked in Density graph it seems that most of the Employees enrolled with us seems to be in 5,00,000-6,00,000 job bracket.

```
In [ ]: fig=plt.figure(figsize=(32,8)) # width*height


for ind_num,col_name in enumerate(df.describe(include='object')):

    # get top n categories
    top_n = 20
    top_categories = df[col_name].value_counts().nlargest(top_n).index

    # filter the df
    df_top_n = df[df[col_name].isin(top_categories)]

    # Create a countplot
    plt.subplot(1,5,ind_num+1)
    sns.countplot(x=col_name, data=df_top_n)
    plt.xticks(rotation=45, ha='right') # Rotate x-axis labels for better visibility

    # Set labels and title
    plt.xlabel(col_name)
    plt.ylabel('Frequency')
    plt.title(f'Top {top_n} {col_name}')
```



- Seems we have duplicate entries for candidates in dataset having multiple years of experience.
- Last CTC update year was in 2021.

```
In [ ]: # Categorical Vs Numerical (CTC)

fig=plt.figure(figsize=(100,4))

# Set a range to exclude outliers (adjust as needed)
lower_limit = np.percentile(df['ctc'], 1)
upper_limit = np.percentile(df['ctc'], 90)

# Filter the data to exclude outliers
filtered_data = df[(df['ctc'] >= lower_limit) & (df['ctc'] <= upper_limit)]

# get top n categories
top_n = 120
top_categories = filtered_data["company_hash"].value_counts().nlargest(top_n).index


# filter the df
df_top_n = filtered_data[filtered_data["company_hash"].isin(top_categories)]

plt.subplot(1,3,1)
sns.boxplot(y='ctc', x='company_hash', data=df_top_n)

plt.xticks(rotation=90, ha='center') # Rotate x-axis labels for better visibility

# Set labels and title
plt.xlabel("company_hash")
plt.ylabel('ctc')
plt.title('Box Plot')
```

Out[]: Text(0.5, 1.0, 'Box Plot')



```
In [ ]: fig=plt.figure(figsize=(100,4)) # width*height

# Set a range to exclude outliers (adjust as needed)
lower_limit = np.percentile(df['ctc'], 0)
upper_limit = np.percentile(df['ctc'], 95)

# Filter the data to exclude outliers
filtered_data = df[(df['ctc'] >= lower_limit) & (df['ctc'] <= upper_limit)]
```



```
# get top n categories
top_n = 120
top_categories = filtered_data["job_position"].value_counts().nlargest(top_n).index


# filter the df
df_top_n = filtered_data[filtered_data["job_position"].isin(top_categories)]

plt.subplot(1,3,1)
sns.boxplot(y='ctc', x='job_position', data=df_top_n)

plt.xticks(rotation=90, ha='center') # Rotate x-axis labels for better visibility

# Set labels and title
plt.xlabel("job_position")
plt.ylabel('ctc')
plt.title(f'Box Plot')
```

Out []: Text(0.5, 1.0, 'Box Plot')



- Profiles seems to earn more irrespective of which company they work for includes:

- SDE-3
- Tech Lead Software Engineer


```
In [ ]: # Category Vs Category

# get top n categories
top_n = 20
top_categories = df["job_position"].value_counts().nlargest(top_n).index

# filter the df
df_top_n = df[df["job_position"].isin(top_categories)]

cross_tab = pd.crosstab(df_top_n["job_position"], df_top_n["ctc_updated_year"])
sns.heatmap(cross_tab, annot=True, cmap='YlGnBu", fmt='d', cbar=True)
```

Out []: <Axes: xlabel='ctc_updated_year', ylabel='job_position'>



- Full Stack Engineers and Backend Engineer have higher chances of getting Appraisal or Salary hike as compare to other Profiles

```
In [ ]: fig=plt.figure(figsize=(100,4)) # width*height

# Set a range to exclude outliers (adjust as needed)
lower_limit = np.percentile(df['ctc'], 0)
upper_limit = np.percentile(df['ctc'], 95)

# Filter the data to exclude outliers
filtered_data = df[(df['ctc'] >= lower_limit) & (df['ctc'] <= upper_limit)]

# get top n categories
top_n = 120
top_categories = filtered_data["orgyear"].value_counts().nlargest(top_n).index


# filter the df
df_top_n = filtered_data[filtered_data["orgyear"].isin(top_categories)]

plt.subplot(1,3,1)
sns.boxplot(y='ctc', x='orgyear', data=df_top_n)

plt.xticks(rotation=90, ha='center') # Rotate x-axis labels for better visibility

# Set labels and title
plt.xlabel("orgyear")
plt.ylabel('ctc')
plt.title(f'Box Plot')
```

Out []: Text(0.5, 1.0, 'Box Plot')



- From above graph we have a lot of noise data, we'll be filter that later at pre-processing stage

```
In [ ]: df['email_hash'].value_counts()[0:10]
```

```
Out [ ]: email_hash
bbace3cc586408bbc65765bc6a16b77d8913836cfc98b77c05488f02f5714a4b 10
3e5e49daa5527a6d5a33599b238bf9bf31e85b9efa9a94f1c88c5e15a6f31378 9
298528ce3160cc761e4dc37a07337ee2e0589df251d73645aae209b010210eee 9
6842660273f70e9aa239026ba33bf8e2275d6a0bd20124021b952b5bc3d07e6c 9
d598d6f1fb21b45593c2af1c2f76ae9f4cb7167156cdf93246d4192a89d8065 8
faf40195f8c58d5c7edc758cc725a762d51920da996410b80ac4a4d85c803da0 8
b4d5afa09bec8689017d8b29701b80d664ca37b83cb883376b2e95191320da66 8
d15041f58bb01c8ee29f72e3b136e26bc32f3169a40b53d75fe7ae9cb9a551 8
4818edfd67ed8563dde5d083306485d91d19f4f1c95d193a1700e79dd245b75c 8
c0eb129061675da412b0deb15871dd06ef0d7cd86eb5f7e8cc6a20b0d1938183 8
Name: count, dtype: int64
```

```
In [ ]: df[df['email_hash']=="bbace3cc586408bbc65765bc6a16b77d8913836cfc98b77c05488f02f5714a4b"]
```


Out[]:

		company_hash	email_hash	orgyear	ctc	job_position	ctc_updated_year
	24109	oxej ntwyzgrgsxto rxbxnta	bbase3cc586400bbc65765bc6a16b77d8913836cfc98b7...	2018.0	720000	NaN	2020.0
	45984	oxej ntwyzgrgsxto rxbxnta	bbase3cc586400bbc65765bc6a16b77d8913836cfc98b7...	2018.0	720000	Support Engineer	2020.0
	72315	oxej ntwyzgrgsxto rxbxnta	bbase3cc586400bbc65765bc6a16b77d8913836cfc98b7...	2018.0	720000	Other	2020.0
	102915	oxej ntwyzgrgsxto rxbxnta	bbase3cc586400bbc65765bc6a16b77d8913836cfc98b7...	2018.0	720000	FullStack Engineer	2020.0
	117764	oxej ntwyzgrgsxto rxbxnta	bbase3cc586400bbc65765bc6a16b77d8913836cfc98b7...	2018.0	720000	Data Analyst	2020.0
	121483	oxej ntwyzgrgsxto rxbxnta	bbase3cc586400bbc65765bc6a16b77d8913836cfc98b7...	2018.0	660000	Other	2019.0
	124476	oxej ntwyzgrgsxto rxbxnta	bbase3cc586400bbc65765bc6a16b77d8913836cfc98b7...	2018.0	660000	Support Engineer	2019.0
	144479	oxej ntwyzgrgsxto rxbxnta	bbase3cc586400bbc65765bc6a16b77d8913836cfc98b7...	2018.0	660000	FullStack Engineer	2019.0
	152801	oxej ntwyzgrgsxto rxbxnta	bbase3cc586400bbc65765bc6a16b77d8913836cfc98b7...	2018.0	660000	Devops Engineer	2019.0
	159835	oxej ntwyzgrgsxto rxbxnta	bbase3cc586400bbc65765bc6a16b77d8913836cfc98b7...	2018.0	660000	NaN	2019.0

- from above analysis we can check that candidate have multiple entries based on job_positions

In []:

```
df[df['email_hash']=="d598d6f1fb21b45593c2afc1c2f76ae9f4cb7167156cdf93246d4192a89d8065"]
```

Out []:

		company_hash	email_hash	orgyear	ctc	job_position	ctc_updated_year
	4401	nnvn wgzohrnrvzj otqcxwto	d598d6f1fb21b45593c2afc1c2f76ae9f4cb7167156cdf...	2018.0	300000	NaN	2020.0
	11331	nnvn wgzohrnrvzj otqcxwto	d598d6f1fb21b45593c2afc1c2f76ae9f4cb7167156cdf...	2018.0	300000	Data Scientist	2020.0
	22412	nnvn wgzohrnrvzj otqcxwto	d598d6f1fb21b45593c2afc1c2f76ae9f4cb7167156cdf...	2018.0	300000	Frontend Engineer	2020.0
	81028	nnvn wgzohrnrvzj otqcxwto	d598d6f1fb21b45593c2afc1c2f76ae9f4cb7167156cdf...	2018.0	400000	Other	2021.0
	90782	nnvn wgzohrnrvzj otqcxwto	d598d6f1fb21b45593c2afc1c2f76ae9f4cb7167156cdf...	2018.0	400000	Backend Engineer	2021.0
	92949	nnvn wgzohrnrvzj otqcxwto	d598d6f1fb21b45593c2afc1c2f76ae9f4cb7167156cdf...	2018.0	400000	Data Scientist	2021.0
	107425	nnvn wgzohrnrvzj otqcxwto	d598d6f1fb21b45593c2afc1c2f76ae9f4cb7167156cdf...	2018.0	400000	NaN	2021.0
	132398	nnvn wgzohrnrvzj otqcxwto	d598d6f1fb21b45593c2afc1c2f76ae9f4cb7167156cdf...	2018.0	400000	Frontend Engineer	2021.0

In []:

```
df[df['email_hash']=="298528ce3160cc761e4dc37a07337ee2e0589df251d73645aae209b010210eee"]
```

Out[]:

	company_hash	email_hash	orgyear	ctc	job_position	ctc_updated_year	
	65909	cvrhtbgbtznhb	298528ce3160cc761e4dc37a07337ee2e0589df251d736...	2018.0	720000	Backend Engineer	2020.0
	72799	cvrhtbgbtznhb	298528ce3160cc761e4dc37a07337ee2e0589df251d736...	2018.0	720000	Research Engineers	2020.0
	82099	cvrhtbgbtznhb	298528ce3160cc761e4dc37a07337ee2e0589df251d736...	2018.0	720000	Other	2020.0
	93495	cvrhtbgbtznhb	298528ce3160cc761e4dc37a07337ee2e0589df251d736...	2018.0	720000	NaN	2020.0
	93783	cvrhtbgbtznhb	298528ce3160cc761e4dc37a07337ee2e0589df251d736...	2018.0	720000	Data Scientist	2020.0
	190903	cvrhtbgbtznhb	298528ce3160cc761e4dc37a07337ee2e0589df251d736...	2018.0	700000	Other	2020.0
	191498	cvrhtbgbtznhb	298528ce3160cc761e4dc37a07337ee2e0589df251d736...	2018.0	700000	Research Engineers	2020.0
	196685	cvrhtbgbtznhb	298528ce3160cc761e4dc37a07337ee2e0589df251d736...	2018.0	700000	Data Scientist	2020.0
	201587	cvrhtbgbtznhb	298528ce3160cc761e4dc37a07337ee2e0589df251d736...	2018.0	700000	Backend Engineer	2020.0

In []:

```
df_cnt = df['email_hash'].value_counts().reset_index()
df_cnt[df_cnt["count"]==2]['email_hash'].tolist()[0:5]
```

Out []:

```
['d3e27dfa3240546390161c8f8e7be3ea5cd8f47dbf7152c9df81f48dc841742e',
 '2dc0dd508944f55ff448d30a5e660174896c0de2d150e68941cf750bc5f1aa4a',
 '702b13ba2005b5d8be91e0b7a9e8c09b3d71bb84e468379e79fbb71c7615cfaaa',
 'd7df6bd598c376ae391518a835780f0bcfac770b010c664f73ba9fde097077ec6',
 '59f5ea9240dc8e6ba6790a2697c3e9b1605320851b816d7dd2211c102c6f21a6']
```

In []:

```
df[df['email_hash']=="5559de74dd698c1ebc14d9653272d0c612970bc5ba206d704394ccaab18c3cc"]
```

Out []:

	company_hash	email_hash	orgyear	ctc	job_position	ctc_updated_year	
	34281	qtamrvvpnqtt	5559de74dd698c1ebc14d9653272d0c612970bc5ba206d...	2019.0	600000	NaN	2021.0
	184547	ztnbtaowgb	5559de74dd698c1ebc14d9653272d0c612970bc5ba206d...	2019.0	340000	FullStack Engineer	2020.0

In []:

```
df_cnt = df['email_hash'].value_counts().reset_index()
df_cnt[df_cnt["count"]==3]['email_hash'].tolist()[0:5]
```

Out []:

```
['7f9fd9949a7a90e322f3a1e72fb1eba7a1ae670778fe90d4a3a804ab87d7ae6a',
 'bf846924f9ebac3e7bd906316d7da0c0f16aab5b34d0d405c921917f4ed544c',
 '2538e1d0a89523f2a4a112e03b8d93e1b9f6a5d05a73fb2bc0cbe0bcb7ec395f',
 '8c7702516ea7c543e8f7758226a996c3bfb6bc889c432c27d1e9df6d34bcd8a8',
 '0d8e3ae92ca1e52184bb396d9ad92bbbd6568bc579274ce9d2723ba380c8451']
```

In []:

```
df[df['email_hash']=="6576d1f1a561eb06bce501ab7ece60baaee8d529902f16712496a18a4287ff64"]
```

Out[]:

	company_hash	email_hash	orgyear	ctc	job_position	ctc_updated_year
	83611	vbvkgz 6576d1f1a561eb06bce501ab7ece60baaee8d529902f16...	2015.0	499999	Other	2020.0
	96777	vbvkgz 6576d1f1a561eb06bce501ab7ece60baaee8d529902f16...	2015.0	499999	Data Analyst	2020.0
	162676	vbvkgz 6576d1f1a561eb06bce501ab7ece60baaee8d529902f16...	2015.0	499999	NaN	2020.0

In []:

```
df[df['email_hash']=="f6327cc669826a2b55dae0cbb69066a8e5e549320b73a6b0ce5cc9d5cc61c5ed"]
```

Out[]:

	company_hash	email_hash	orgyear	ctc	job_position	ctc_updated_year	
	41961	vbtqxwvz tduqtoo	f6327cc669826a2b55dae0cbb69066a8e5e549320b73a6...	2014.0	2000000	NaN	2021.0
	56796	vbtqxwvz tduqtoo	f6327cc669826a2b55dae0cbb69066a8e5e549320b73a6...	2014.0	2000000	Engineering Leadership	2021.0
	67497	vbtqxwvz tduqtoo	f6327cc669826a2b55dae0cbb69066a8e5e549320b73a6...	2014.0	2000000	FullStack Engineer	2021.0

In []:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 205843 entries, 0 to 205842
Data columns (total 6 columns):
#   Column          Non-Null Count  Dtype
---  -
0   company_hash    205799 non-null  object
1   email_hash      205843 non-null  object
2   orgyear         205757 non-null  object
3   ctc             205843 non-null  int64
4   job_position    153279 non-null  object
5   ctc_updated_year 205843 non-null  object
dtypes: int64(1), object(5)
memory usage: 9.4+ MB
```

In []:

```
missing_value = round((df.job_position.isnull().sum()/df.shape[0])*100 , 2)
missing_value
```

Out []:

```
25.54
```

- 25% of the values in Job Description seems missing and upon analyzing the feature doesn't seems reliable. It is better to remove the feature while moving forward.

In []:

```
df.drop(columns=["job_position"], axis=1, inplace=True)
```

In []:

```
df.head()
```

Out[]:

	company_hash	email_hash	orgyear	ctc	ctc_updated_year
0	atrgxnnt xzaxv	6de0a4417d18ab14334c3f43397fc13b30c35149d70c05...	2016.0	1100000	2020.0
1	qtrxvzwt xzegwgbb rxbxnta	b0aaf1ac138b53cb6e039ba2c3d6604a250d02d5145c10...	2018.0	449999	2019.0
2	ojzwmnwnxw vx	4860c670bcd48fb96c02a4b0ae3608aefdd98176112e9...	2015.0	2000000	2020.0
3	ngpgutaxv	effdede7a2e7c2af664c8a31d9346385016128d66bbc58...	2017.0	700000	2019.0
4	qxen sqghu	6ff54e709262f55cb999a1c1db8436cb2055d8f79ab520...	2017.0	1400000	2019.0

In []:

```
# tbv_diff_pattern : based on box plot
Q1 = df["ctc"].quantile(0.01)
Q2 = df["ctc"].quantile(0.15)
Q3 = df["ctc"].quantile(0.85)
Q4 = df["ctc"].quantile(0.99)

print(df['ctc'].min(), Q1, Q2, Q3, Q4, df['ctc'].max())
```

- ```
2 37000.0 0 400000.0 2300000.0 12600000.0 1000150000
```
- I seems the spread for CTC ranges from 37,000 to 1,26,00,000 which is huge margin and we have to be specific while handling the specific CTC ranges.

In [ ]:

## Data Preprocessing

In [ ]:

```
df.head()
```

Out[ ]:

|   | company_hash              | email_hash                                        | orgyear | ctc     | ctc_updated_year |
|---|---------------------------|---------------------------------------------------|---------|---------|------------------|
| 0 | atrgxnnt xzaxv            | 6de0a4417d18ab14334c3f43397fc13b30c35149d70c05... | 2016.0  | 1100000 | 2020.0           |
| 1 | qtrxvzwt xzegwgbb rxbxnta | b0aaf1ac138b53cb6e039ba2c3d6604a250d02d5145c10... | 2018.0  | 449999  | 2019.0           |
| 2 | ojzwmnwnxw vx             | 4860c670bcd48fb96c02a4b0ae3608aefdd98176112e9...  | 2015.0  | 2000000 | 2020.0           |
| 3 | ngpgutaxv                 | effdede7a2e7c2af664c8a31d9346385016128d66bbc58... | 2017.0  | 700000  | 2019.0           |
| 4 | qxen sqghu                | 6ff54e709262f55cb999a1c1db8436cb2055d8f79ab520... | 2017.0  | 1400000 | 2019.0           |

In [ ]:

```
df.shape
```

Out [ ]:

```
(205843, 5)
```

In [ ]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 205843 entries, 0 to 205842
Data columns (total 5 columns):
Column Non-Null Count Dtype
--- -
0 company_hash 205799 non-null object
1 email_hash 205843 non-null object
2 orgyear 205757 non-null object
3 ctc 205843 non-null int64
4 ctc_updated_year 205843 non-null object
dtypes: int64(1), object(4)
memory usage: 7.9+ MB
```

```
In []: df["orgyear"].value_counts()
```

```
Out[]: orgyear
2018.0 25256
2019.0 23427
2017.0 23239
2016.0 23043
2015.0 20610
...
4.0 1
1900.0 1
1971.0 1
201.0 1
200.0 1
Name: count, Length: 77, dtype: int64
```

- Points to consider while filtering orgyear
1. We have total experience values it seems in this column as well.
  2. Some noise as well with values like 201, 200 etc
  3. May contain negative values also.

```
In []: df["orgyear"].max()
```

```
Out[]: 20165.0
```

```
In []: df["orgyear"].min()
```

```
Out[]: 0.0
```

```
In []: df["year_exp"] = df["orgyear"].apply(lambda x: (2024-x) if (2024-x)<50 else x)
df.head()
```

```
Out[]: company_hash email_hash orgyear ctc ctc_updated_year year_exp
0 atrgxmnt xzaxv 6de0a4417d18ab14334c3f43397fc13b30c35149d70c05... 2016.0 1100000 2020.0 8.0
1 qtrxvzwt xzegwgb rxbxnta b0aaf1ac138b53cb6e039ba2c3d6604a250d02d5145c10... 2018.0 449999 2019.0 6.0
2 ojzwnvwnxw vx 4860c670bcd48fb96c02a4b0ae3608ae6fdd98176112e9... 2015.0 2000000 2020.0 9.0
3 ngpgutaxv effdede7a2e7c2af664c8a31d9346385016128d66bbc58... 2017.0 700000 2019.0 7.0
4 qxen sqghu 6ff54e709262f55cb999a1c1db8436cb2055d8f79ab520... 2017.0 1400000 2019.0 7.0
```

```
In []: df.shape
```

```
Out[]: (205843, 6)
```

```
In []: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 205843 entries, 0 to 205842
Data columns (total 6 columns):
Column Non-Null Count Dtype
--- -
0 company_hash 205799 non-null object
1 email_hash 205843 non-null object
2 orgyear 205757 non-null object
3 ctc 205843 non-null int64
4 ctc_updated_year 205843 non-null object
5 year_exp 205757 non-null float64
dtypes: float64(1), int64(1), object(4)
memory usage: 9.4+ MB
```

```
In []: df['year_exp'] = df['year_exp'].apply(lambda x: x if x<=50 and x>=0 else np.nan)
```

```
In []: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 205843 entries, 0 to 205842
Data columns (total 6 columns):
Column Non-Null Count Dtype
--- -
0 company_hash 205799 non-null object
1 email_hash 205843 non-null object
2 orgyear 205757 non-null object
3 ctc 205843 non-null int64
4 ctc_updated_year 205843 non-null object
5 year_exp 205699 non-null float64
dtypes: float64(1), int64(1), object(4)
memory usage: 9.4+ MB
```

```
In []: fig=plt.figure(figsize=(100,4)) # width*height

Set a range to exclude outliers (adjust as needed)
lower_limit = np.percentile(df['ctc'], 0)
upper_limit = np.percentile(df['ctc'], 95)

Filter the data to exclude outliers
filtered_data = df[(df['ctc'] >= lower_limit) & (df['ctc'] <= upper_limit)]

get top n categories
top_n = 120
top_categories = filtered_data["year_exp"].value_counts().nlargest(top_n).index


filter the df
df_top_n = filtered_data[filtered_data["year_exp"].isin(top_categories)]

plt.subplot(1,3,1)
sns.boxplot(y='ctc', x='year_exp', data=df_top_n)

plt.xticks(rotation=90, ha='center') # Rotate x-axis labels for better visibility

Set labels and title
plt.xlabel("year_exp")
plt.ylabel('ctc')
plt.title(f'Box Plot')
```

```
Out[]: Text(0.5, 1.0, 'Box Plot')
```



- CTC seems to have noise after 40, so we'll be filtering the values till 40 years of exp.

```
In []: df['year_exp'] = df['year_exp'].apply(lambda x: x if x<=40 and x>=0 else np.nan)

fig=plt.figure(figsize=(100,4)) # width*height

Set a range to exclude outliers (adjust as needed)
lower_limit = np.percentile(df['ctc'], 0)
upper_limit = np.percentile(df['ctc'], 99)

Filter the data to exclude outliers
filtered_data = df[(df['ctc'] >= lower_limit) & (df['ctc'] <= upper_limit)]

get top n categories
top_n = 120
top_categories = filtered_data["year_exp"].value_counts().nlargest(top_n).index

filter the df
df_top_n = filtered_data[filtered_data["year_exp"].isin(top_categories)]


plt.subplot(1,3,1)
sns.boxplot(y='ctc', x='year_exp', data=df_top_n)

plt.xticks(rotation=90, ha='center') # Rotate x-axis labels for better visibility

Set labels and title
plt.xlabel("year_exp")
plt.ylabel('ctc')
plt.title(f'Box Plot')
```

```
Out[]: Text(0.5, 1.0, 'Box Plot')
```





```
In []: df.head()

Out[]:
 company_hash email_hash orgyear ctc ctc_updated_year year_exp
0 atrgxnnnt xzaxv 6de0a4417d18ab14334c3f43397fc13b30c35149d70c05... 2016.0 1100000 2020.0 8.0
1 qtrxvzwt xzegwgb rxbxnta b0aaf1ac138b53cb6e039ba2c3d6604a250d02d5145c10... 2018.0 449999 2019.0 6.0
2 ojzwnvwnxw vx 4860c670bcd48fb96c02a4b0ae3608ae6fdd98176112e9... 2015.0 2000000 2020.0 9.0
3 ngpgutaxv effdede7a2e7c2af664c8a31d9346385016128d66bbc58... 2017.0 700000 2019.0 7.0
4 qxen sqghu 6ff54e709262f55cb999a1c1db8436cb2055d8f79ab520... 2017.0 1400000 2019.0 7.0
```

- email\_hash and orgyear doesn't seems to be of any use for us

```
In []: df.drop(columns=["orgyear"], inplace=True, axis=1)

In []: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 205843 entries, 0 to 205842
Data columns (total 5 columns):
Column Non-Null Count Dtype
--- -
0 company_hash 205799 non-null object
1 email_hash 205843 non-null object
2 ctc 205843 non-null int64
3 ctc_updated_year 205843 non-null object
4 year_exp 205691 non-null float64
dtypes: float64(1), int64(1), object(3)
memory usage: 7.9+ MB
```

```
In []: missing_value = round((df.year_exp.isnull().sum()/df.shape[0])*100 , 2)
print(f"year_exp missing value percentage: {missing_value}")

missing_value = round((df.company_hash.isnull().sum()/df.shape[0])*100 , 2)
print(f"company_hash missing value percentage: {missing_value}")

year_exp missing value percentage: 0.07
company_hash missing value percentage: 0.02
```

- Data seems missing let's merge and see the changes

```
In []: df.drop_duplicates(inplace=True)

In []: df.shape

Out[]: (166861, 5)

In []: df.head()
```

```
Out[]:
 company_hash email_hash ctc ctc_updated_year year_exp
0 atrgxnnnt xzaxv 6de0a4417d18ab14334c3f43397fc13b30c35149d70c05... 1100000 2020.0 8.0
1 qtrxvzwt xzegwgb rxbxnta b0aaf1ac138b53cb6e039ba2c3d6604a250d02d5145c10... 449999 2019.0 6.0
2 ojzwnvwnxw vx 4860c670bcd48fb96c02a4b0ae3608ae6fdd98176112e9... 2000000 2020.0 9.0
3 ngpgutaxv effdede7a2e7c2af664c8a31d9346385016128d66bbc58... 700000 2019.0 7.0
4 qxen sqghu 6ff54e709262f55cb999a1c1db8436cb2055d8f79ab520... 1400000 2019.0 7.0
```

```
In []: df.email_hash.value_counts()

Out[]:
email_hash
c986600ef19093ce70837408516acac9570566a4b29b554cfb6b744ffbe697d6 3
db84980ad197f8eff08b14a3442ff57f6374ea780f2587b310aac54b6c32ee3a 3
e5960ec01a207bfa6b83d5576f3d66f98b95f2e200f250303027c95395e4bad9 2
607910ba90b77a948e9255d472b9281ce90e10bc2a3089fedee923c0eb421d039 2
1ea4e620e2d1f02c6c546c1d263582badacdb4b783d4cb8660901d3940f7ac3b 2
..
611b800e5d0f4caae6ade0c9392d27590987e77237dc430a2dd7c95d8bb82fa4 1
14d7bfecc80ded3532b8cf37e5b20989448b5d3fe387862b0f295d3117056a5e 1
2fbd7838b0092973720ba4bf6775f07b63cf88373571273c350aaf4652f835fa 1
8fffb1aeal2ae1dcf3ee1cfe2934117af26df3e989dae5b050aaed826a46c8d 1
bad0e8acadb3e3bddfee92367413ec947bbb1029826fffc3becaea02593a10f8 1
Name: count, Length: 153443, dtype: int64

In []: df[df["email_hash"]=="c986600ef19093ce70837408516acac9570566a4b29b554cfb6b744ffbe697d6"]

Out[]:
 company_hash email_hash ctc ctc_updated_year year_exp
8695 oyxcozvd uqxcvnt rxbxnta c986600ef19093ce70837408516acac9570566a4b29b55... 1800000 2021.0 11.0
61154 uvqrt c986600ef19093ce70837408516acac9570566a4b29b55... 3500000 2021.0 15.0
98347 oyxcozvd uqxcvnt rxbxnta c986600ef19093ce70837408516acac9570566a4b29b55... 1800000 2019.0 11.0
```

```
In []: # Find the index of the rows with the maximum 'Year' within each group
idx = df.groupby(['email_hash','company_hash'])['ctc_updated_year'].idxmax()

Use the index to select the corresponding rows from the original DataFrame
df = df.loc[idx]

Reset index for the final DataFrame
df.reset_index(drop=True, inplace=True)
```

```
In []: df[df["email_hash"]=="c986600ef19093ce70837408516acac9570566a4b29b554cfb6b744ffbe697d6"]

Out[]:
 company_hash email_hash ctc ctc_updated_year year_exp
126055 oyxcozvd uqxcvnt rxbxnta c986600ef19093ce70837408516acac9570566a4b29b55... 1800000 2021.0 11.0
126056 uvqrt c986600ef19093ce70837408516acac9570566a4b29b55... 3500000 2021.0 15.0
```

```
In []: df.shape

Out[]: (160273, 5)

In []: df.email_hash.value_counts()
```

```
Out[]:
email_hash
db84980ad197f8eff08b14a3442ff57f6374ea780f2587b310aac54b6c32ee3a 3
e49d643e06c85681ee3b6feff020134f63c92b17e637283d44854412eb2c95fb 2
9c20e7a5e0b46a4350327978c130282184c2b39772405c5f5c9ab7e1e03b69e8 2
c88c616fb855fd35009d643fc4d9d91b4d53d363057fdc7fd071717c811296e6 2
e496afal17c48da3980dbd8330caf89ec01bd4d4849d17ad87a2b98abb699d172 2
..
57972fe710544a37de4ed56a1ac6f242d25d3f91cac5a4b78b840bea2c6fce30 1
57978d4cf69f3aa592a7a9a6eb5aecd0f5aa25005937befb846d16a6fd4aee6 1
57983cc12ab513f649544370ad98a61a83e64a228a8d8b1a77c73cfc51b0b2e7 1
5798a70eb4780ddb0087a3c11ff7f6d7e23731ebf92dddeb3b7ca0b4f5c6df2d 1
5794bec6e3ee46ffdb464cd55aa9b34a1e5bc3fc08e52cd58570861bed8b8a88 1
Name: count, Length: 153411, dtype: int64
```

```
In []: df[df["email_hash"]=="db84980ad197f8eff08b14a3442ff57f6374ea780f2587b310aac54b6c32ee3a"]

Out[]:
 company_hash email_hash ctc ctc_updated_year year_exp
137367 vqwotqct db84980ad197f8eff08b14a3442ff57f6374ea780f2587... 700000 2021.0 4.0
137368 vqwotqct xzaxv ogrhnxgo rxbxnta db84980ad197f8eff08b14a3442ff57f6374ea780f2587... 700000 2021.0 4.0
137369 zgn vuurxwvmt db84980ad197f8eff08b14a3442ff57f6374ea780f2587... 400000 2019.0 4.0
```

```
In []: df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 160273 entries, 0 to 160272
Data columns (total 5 columns):
Column Non-Null Count Dtype
--- -
0 company_hash 160273 non-null object
1 email_hash 160273 non-null object
2 ctc 160273 non-null int64
3 ctc_updated_year 160273 non-null object
4 year_exp 160134 non-null float64
dtypes: float64(1), int64(1), object(3)
memory usage: 6.1+ MB
```

```
In []: missing_value = round((df.year_exp.isnull().sum()/df.shape[0])*100 , 2)
print(f"year_exp missing value percentage: {missing_value}")

year_exp missing value percentage: 0.09
```

- only .09% percent values seems to be missing we can drop the specific rows / perform imputing on that dataset.

```
In []: ## Missing value treatment

from sklearn.impute import KNNImputer

numeric_columns = ["ctc", "year_exp"]

Extract the numeric part of the DataFrame
df_numeric = df[numeric_columns]

Impute missing values in the numeric part using KNNImputer
knn_imputer = KNNImputer(n_neighbors=20)
df_numeric_imputed = pd.DataFrame(knn_imputer.fit_transform(df_numeric), columns=numeric_columns)

Combine the imputed numeric part with the original categorical part
df = pd.concat([df.drop(columns=numeric_columns), df_numeric_imputed], axis=1)
df.isnull().sum()
```

```
Out[]: company_hash 0
email_hash 0
ctc_updated_year 0
ctc 0
year_exp 0
dtype: int64
```

```
In []: df.shape
```

```
Out[]: (160273, 5)
```

```
In []: df.drop(columns=["email_hash"], axis=1, inplace=True)
```

```
In []: df.head()
```

```
Out[]: company_hash ctc_updated_year ctc year_exp
0 bxwqgogen 2019.0 3500000.0 12.0
1 nqsn axsnvr 2020.0 250000.0 11.0
2 gunhb 2019.0 1300000.0 3.0
3 bxwqgotbx wqgugvnxgz 2021.0 2000000.0 20.0
4 fvrbvqn rvmc 2018.0 3400000.0 15.0
```

```
In []: df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 160273 entries, 0 to 160272
Data columns (total 4 columns):
Column Non-Null Count Dtype
--- -
0 company_hash 160273 non-null object
1 ctc_updated_year 160273 non-null object
2 ctc 160273 non-null float64
3 year_exp 160273 non-null float64
dtypes: float64(2), object(2)
memory usage: 4.9+ MB
```

```
In []: ## Outlier Treatment
...
Based on our previous understanding we will be capping the values at 1% i.e. below 37,000 and above 1,26,00,000
...

Set upper and lower thresholds for capping (adjust as needed)
upper_threshold = 12600000
lower_threshold = 37000

Apply capping to remove outliers
df['ctc'].clip(lower=lower_threshold, upper=upper_threshold, inplace=True)

df.shape

/tmp/ipykernel_48614/1913124230.py:11: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.
The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

df['ctc'].clip(lower=lower_threshold, upper=upper_threshold, inplace=True)
```

```
Out[]: (160273, 4)
```

- Doesn't seems there exist any outlier after we did merging and filter the noise.

```
In []: df.head()
```

```
Out[]: company_hash ctc_updated_year ctc year_exp
0 bxwqgogen 2019.0 3500000.0 12.0
1 nqsn axsnvr 2020.0 250000.0 11.0
2 gunhb 2019.0 1300000.0 3.0
3 bxwqgotbx wqgugvnxgz 2021.0 2000000.0 20.0
4 fvrbvqn rvmc 2018.0 3400000.0 15.0
```

```
In []: ## Feature Engineering
...
1. Calculated Years of experience using 'orgyear' features.
2. Based on Job Profile the data seems fine for clustering and Job Description seems not reliable as per understanding.
3. CTC_updated year: we will be considering the feature as categorical value.
4. Implementation of regex for the usecase doesn't seems to be a good idea here.
...

Define bin edges and labels
bins = [0, 500000, 1500000, 3000000, float('inf')] # 'inf' represents infinity, covering values greater than 120
bin_labels = ['Low', 'Medium', 'High', 'Very High']

Create a new column 'Bins' and assign the bin labels
df['salary_bin'] = pd.cut(df['ctc'], bins=bins, labels=bin_labels, right=False)

Define bin edges and labels
bins = [0, 5, 10, 50] # 'inf' represents infinity, covering values greater than 120
bin_labels = ['Low', 'Medium', 'High']

Create a new column 'Bins' and assign the bin labels
df['exp_bin'] = pd.cut(df['year_exp'], bins=bins, labels=bin_labels, right=False)
```

```
In []: # Define conditions and corresponding labels
conditions = [
 (df['ctc_updated_year'] == 2021),
 (df['ctc_updated_year'] == 2020),
 (df['ctc_updated_year'] != 2020) & (df['ctc_updated_year'] != 2021)
]
labels = ['2021', '2020', 'other']
df['ctc_updated'] = np.select(conditions, labels, default='Other')
```

```
In []: df['ctc_updated'].value_counts()
```

```
Out[]: ctc_updated
other 78368
2021 42742
2020 39163
Name: count, dtype: int64
```

```
In []: df.drop(columns=["ctc", "year_exp", "ctc_updated_year"], axis=1, inplace=True)
```

```
In []: df.drop_duplicates(inplace=True)
```

```
In []: df.describe(include="all")
```

```
Out[]: company_hash salary_bin exp_bin ctc_updated
count 66898 66898 66898 66898
unique 37299 4 3 3
top bxwqgogen Medium Medium other
freq 36 30519 33570 32612
```

```
In []: tmpdf = df.company_hash.value_counts().reset_index()
```

```
In []: for i in range(1,36):
 print(i, tmpdf[tmpdf["count"]==i].shape[0])
```

1 29137  
2 3707  
3 1418  
4 748  
5 487  
6 346  
7 271  
8 172  
9 129  
10 135  
11 103  
12 99  
13 70  
14 70  
15 37  
16 45  
17 39  
18 41  
19 38  
20 37  
21 26  
22 20  
23 17  
24 13  
25 10  
26 10  
27 15  
28 10  
29 10  
30 5  
31 7  
32 6  
33 7  
34 3  
35 5

```
In []: # Count the occurrences of each category
category_counts = df['company_hash'].value_counts()

Identify categories with count 1
single_occurrence_categories = category_counts[category_counts == 1].index

Replace single-occurrence categories with 'Other'
df['company_hash'] = df['company_hash'].replace(single_occurrence_categories, 'Other')
```

```
In []: df.drop_duplicates(inplace=True)
```

```
In []: df.shape
```

Out[ ]: (37797, 4)

```
In []: df.company_hash.value_counts()
```

Out[ ]: company\_hash  
xzegojo 36  
bxwqgogen 36  
xmb 36  
Other 36  
vbvkgz 36  
.  
.  
hzxcvqxtnj 2  
btnnrtqngrtag xzntqzvnxgzvr xzw 2  
hubw tzntquqxoto 2  
wvubvnqxd ntwyzgrgsj 2  
x3 wgzohrnxs 2  
Name: count, Length: 8163, dtype: int64

```
In []: df.describe()
```

Out[ ]:

|        | company_hash | salary_bin | exp_bin | ctc_updated |
|--------|--------------|------------|---------|-------------|
| count  | 37797        | 37797      | 37797   | 37797       |
| unique | 8163         | 4          | 3       | 3           |
| top    | xzegojo      | Medium     | Medium  | other       |
| freq   | 36           | 16522      | 18909   | 16806       |

```
In []: df.describe().columns
```

Out[ ]: Index(['company\_hash', 'salary\_bin', 'exp\_bin', 'ctc\_updated'], dtype='object')

```
In []: ## Data preparation and Scaling

OHE
df = pd.get_dummies(df, columns=['salary_bin', 'exp_bin', 'ctc_updated'], dtype=int)
```

```
In []: ## Target Encoding
import category_encoders as ce
target_encoder = ce.TargetEncoder(cols=['company_hash']) # Create a TargetEncoder instance
df = target_encoder.fit_transform(df, df.index) # Fit and transform the DataFrame with target encoding, using index as a dummy target

Standard Scaler
from sklearn.preprocessing import MinMaxScaler
scaler = MinMaxScaler()

Reshape the data to fit the scaler (required for a single column)
df["company_hash"] = scaler.fit_transform(df["company_hash"].values.reshape(-1, 1))
```

```
In []: df.head()
```

Out[ ]:

|   | company_hash | salary_bin_Low | salary_bin_Medium | salary_bin_High | salary_bin_Very High | exp_bin_Low | exp_bin_Medium | exp_bin_High | ctc_updated_2020 | ctc_updated_2021 | ctc_updated_other |
|---|--------------|----------------|-------------------|-----------------|----------------------|-------------|----------------|--------------|------------------|------------------|-------------------|
| 0 | 0.166218     | 0              | 0                 | 0               | 1                    | 0           | 0              | 1            | 0                | 0                | 1                 |
| 1 | 0.000000     | 1              | 0                 | 0               | 0                    | 0           | 0              | 1            | 1                | 0                | 0                 |
| 2 | 0.491054     | 0              | 1                 | 0               | 0                    | 1           | 0              | 0            | 0                | 0                | 1                 |
| 3 | 0.788863     | 0              | 0                 | 1               | 0                    | 0           | 0              | 1            | 0                | 1                | 0                 |
| 4 | 0.351813     | 0              | 0                 | 0               | 1                    | 0           | 0              | 1            | 0                | 0                | 1                 |

```
In []: df.to_csv("dataset/phase2_df.csv", index=False)
```

```
In []:
```

```
In []:
```